

CS201 Spring 2026 — Lecture Notes

Clarissa Littler

Spring 2026

Preface

This document is a running narrative of what we covered in each class meeting this term. Each section pairs the agenda for the day with the little code samples we actually typed up together, so if you missed a lecture or want to go back and re-read, the whole arc of a day should be here in one place. The source files themselves live in the `lectureN/` subdirectories of the class repo, and all of the snippets quoted below are verbatim from those files.

Contents

1	Lecture 1 — Welcome, C vs. C++, Bitwise Tricks, and a Peek at Integers	1
2	Lecture 2 — Hex, Endianness, IEEE-754, and Where Floats Lie to You	3
3	Lecture 3 — malloc, and Your First Real Assembly	5
4	Lecture 4 — Registers, Calling Convention, Addressing Modes, Control Flow, and “Hello, world”	7
5	Lecture 5 — Hello Again, echo, and the Road from Jumps to Proper Functions	11
6	Lecture 6 — Reusing Functions Across Files, Recursion, and Summing an Array	16
7	Lecture 7 — Processes, fork, Threads, and Why Sharing State Hurts	20
8	Lecture 8 — Returning From Threads, Mutexes, and Critical Sections	26
1	Lecture 1 — Welcome, C vs. C++, Bitwise Tricks, and a Peek at Integers	

This one was mostly an orientation plus a crash-course in the parts of C that tend to trip up people coming from C++. I was genuinely a little surprised everyone showed up, so thank you for that!

Logistics

The class is my own personal admixture of four things:

- C

- x86-64 assembly
- Concurrency
- Operating systems theory and practice

Grading-wise, you’ve got quizzes, discussions, and coding assignments. The class repo is going to be your friend: that’s where the textbook lives, where all the “Agenda” files like this one come from, and where we’ll drop every bit of code we work on in class.

Structurally, I’ll aim to start about five minutes after 10, break after every 50 minutes, and try not to keep you much past two hours.

C: What’s different?

If you’re coming from C++, the big shocks are:

- No pass-by-reference. You get pass-by-value or you get pointers, full stop.
- No `bool` type (well, not in the way you’re used to — nonzero is truthy and zero is falsy, and there’s no native `true/false`).
- No `nullptr`; you use the `NULL` macro (which is basically just 0 cast to a pointer).
- I/O is `printf/scanf` rather than streams.

We played with two small files to make this concrete. The first, `arraything.c`, shows that arrays do decay to pointers — so you pass the length by hand — and that you iterate with a plain `for` loop rather than an iterator:

```
void arrayMunger(int a[], int s){
    for(int i=0; i<s; i++){
        a[i] = 2*a[i];
    }
}
```

The little commentary `// arr[i] ==> *(arr + i)` hints at the thing we’ll lean on all term: indexing is just pointer arithmetic in disguise.

The second file, `passPointy.c`, makes the “no pass by reference” point tangible. If you want a function to modify its caller’s variable, you pass the *address* and dereference inside:

```
void intMunger(int* p){
    *p = 2*(*p);
}
```

And you call it with `intMunger(&test);`. That ampersand is doing real work — it’s giving you a pointer to `test` rather than a copy of it.

Bitwise operations

We went through the usual suspects with truth tables:

- `&` (AND), `|` (OR), `^` (XOR), `~` (NOT).

- << (shift left, pads with zeroes, anything that falls off the edge is gone) and >> (shift right, same idea on the other side).

The file `bitmunger1.c` is where this all came together. It defines a little loop that prints out the 32 bits of an `int`, lets you pick a bit to toggle, and optionally negates the integer via two's complement:

```
void toggleBit(int* num, int place){
    *num = (*num) ^ (1 << place);
}

void twosComplement(int* num){
    *num = ~((*num) - 1);
}
```

`toggleBit` is the canonical XOR trick: XOR-ing a bit with 1 flips it, XOR-ing with 0 leaves it alone. `twosComplement` implements the “subtract one, then invert” rule for negating a two's-complement integer.

Peeking inside an integer: two's complement

We derived why two's complement works. A plain unsigned n -bit integer encodes

$$\sum_{i=0}^{n-1} b_i \cdot 2^i,$$

so 11111111 (8 bits) is $128 + 64 + \dots + 1 = 255$. In two's complement, the top bit's weight flips sign:

$$-b_{n-1} \cdot 2^{n-1} + \sum_{i=0}^{n-2} b_i \cdot 2^i.$$

That gives you one clean mental check: all-ones is always -1 , because $-2^{n-1} + (2^{n-1} - 1) = -1$. And to negate a number, you subtract one then invert all the bits (or equivalently, invert and then add one).

A preview of floats

We ran out of time before we could fully explore floats, but I showed off `floatmunger1.c`, which is structurally the same as `bitmunger1.c` but reinterprets an `int`'s bits as a `float` via a pointer cast. The little spaces it prints (between bits 31/30 and 23/22) are hinting at the IEEE-754 sign / exponent / fraction split we'll dig into next time.

2 Lecture 2 — Hex, Endianness, IEEE-754, and Where Floats Lie to You

The plan was to finish the float discussion we previewed on day one, but before diving in we took a small detour through hex and endianness.

Warm-up: hex codes and the shape of a pointer

Hex is base-16, digits 0–F. The reason it shows up everywhere — addresses, colors, byte dumps — is that two hex digits are exactly one byte ($16 \cdot 16 = 256 = 2^8$). So when you see an RGBA color like #8765AF, that’s really three (or four) bytes: one for red, one for green, one for blue, optionally one for transparency.

Endianness

We wrote `byteOrderTest.c` to actually look at how an 8-byte integer is laid out in memory on our machines:

```
long int num = 0x0123456789abcdef;
unsigned char* p = (unsigned char*)&num;

printf("The bytes are stored in memory as: ");
for(int i=0; i<8; i++){
    printf("%x ",*(p+i));
}
```

On an x86-64 box this prints the bytes out in the opposite order from how we wrote the literal — classic little-endian. The trick is the `(unsigned char*)` cast: once we have a byte-pointer, pointer arithmetic moves us one byte at a time, so we can walk the actual storage.

Sizes and pointers

Alongside that, `sizes.c` just prints `sizeof(int)`, `sizeof(long)`, `sizeof(char)`, `sizeof(float)`, `sizeof(double)`, and the size of a pointer. On our 64-bit Linux boxes: `int` is 4, `long` is 8, `char` is 1, `float` is 4, `double` is 8, and any pointer is 8. That last one is the important bit — *all* pointers are the same size, regardless of what they point to, because they’re all just addresses.

And `pointerArith.c` drove home that pointer arithmetic scales with the *pointee* size. An `int*` advances 4 bytes per `+1`, a `char*` advances 1, a `char**` advances 8, and so on.

IEEE-754, formally

We did the 32-bit (`float`) version because the 64-bit version is just “same thing with bigger fields.”

- 1 bit sign
- 8 bits exponent (stored with a bias of 127)
- 23 bits fraction

In the normal case, the value is

$$(-1)^s \cdot 2^{e-127} \cdot (1 + f),$$

where f is the fractional part read as “rational binary”: the bit at position k after the implicit point contributes 2^{-k} .

Worked example: 0 10000001 1010...0 is

$$(-1)^0 \cdot 2^{129-127} \cdot (1 + \frac{1}{2} + \frac{1}{8}) = 4 \cdot \frac{13}{8} = 6.5.$$

We exercised this live with `floatmunger1.c` (same idea as lecture 1, but now with `%.16f` so you can see lots of digits). Flipping specific bits and watching the printed value change is the single best way to internalize the format.

Edge cases

- Exponent is all 1s:
 - fraction is zero $\Rightarrow \pm\infty$,
 - fraction is nonzero $\Rightarrow \text{NaN}$.
- Exponent is all 0s (subnormal numbers): the exponent sticks at 2^{-126} , but the leading implicit 1 is gone, so the value is $(-1)^s \cdot 2^{-126} \cdot f$.

Where floats lie to you

`floatlimit.c` is the “gotcha” example:

```
float f = 0.0001;
float accum = 0;
for(int i=0; i<10000; i++){
    accum += f;
}
printf("%.20f\n", accum);
```

We’d love this to print `1.00000000000000000000`. It doesn’t. The reason is a mix of: `0.0001` is not exactly representable in binary, and we’re accumulating that small rounding error 10,000 times.

Why is this *unavoidable*? A bit of math: the rationals and the integers are countably infinite, but the reals form a strictly larger uncountable set. There is no way to cover the reals with finitely many bits, or even with countably many bits; we’re forced to pick a countable scatter of “breadcrumbs” and round every other real to the nearest one.

Round to even

To keep that rounding error from biasing one direction, IEEE-754 specifies “round to even”: when a real is exactly halfway between two breadcrumbs, the hardware rounds to the one whose last bit is zero. Over many operations this rounds up about as often as it rounds down.

3 Lecture 3 — malloc, and Your First Real Assembly

The real headline today: we left C-land and wrote actual assembly against the bare Linux syscall interface. But first, I owed you `malloc`.

malloc, finally

In my rush last time I completely skipped dynamic allocation, which is embarrassing, so we opened with two quick examples.

`malloc.c` is the minimum viable pointer dance:

```
int* n = malloc(sizeof(int));
*n = 10;
printf("%d\n", *n);
free(n);
n = 0;
if(0 == NULL){
    printf("Yup\n");
}
```

Two things to note. First, `malloc` always returns a pointer and always needs you to tell it how many bytes you want — unlike `new T` in C++, it doesn't know anything about your type, so you pass `sizeof(int)` (or a multiple of it). Second, that final `if` is my excuse to point out that `NULL` is literally just 0 cast to a pointer type.

`malloc2.c` extends this to an array and a struct:

```
int* narr = malloc(100*sizeof(int));
narr[2] = 5;

struct Goofus* g = malloc(sizeof(struct Goofus));
g->g1 = 5; // same thing as (*g).g1
```

Once you've got a pointer from `malloc`, you can index it just like an array. The `->` arrow is shorthand for “dereference then access a field.”

Stepping back: what is a CPU, anyway?

The big conceptual pivot for the rest of the term:

- A CPU is, at bottom, a thing that fetches *instructions* from memory and does what they say.
- The set of instructions it understands is an *instruction set architecture* (ISA). Ours is x86-64.
- Assembly language is a one-to-one (ish) textual encoding of those instructions, so writing assembly is as close to “talking to the CPU” as we'll practically get.
- Registers are a tiny amount of extremely fast storage inside the CPU itself. Most instructions operate on registers; memory comes in via load/store.
- Memory hierarchy, the 30-second version: registers → L1/L2/L3 cache → main memory → disk. Each level is roughly an order of magnitude slower and roughly an order of magnitude bigger than the one above it.

I also mentioned `gcc -S`, which tells `gcc` to stop at assembly rather than assemble and link. We'll come back to it.

Your first bare assembly program

`first.s` is deliberately broken:

```
.section .text
.global _start

_start:
```

```
mov $60,%rax
```

This is the shortest thing that looks like an assembly program. The `.text` section is where code lives; `_start` is the symbol the linker uses as the program entry point (unlike in a C program, there is no runtime to call `main` for us).

We moved 60 into `%rax` — 60 is the syscall number for `exit` on x86-64 Linux — and then... did nothing else. Because we never actually issued the `syscall` instruction and there's no implicit return, the CPU just wandered off the end of our code and trapped. Hence the crash.

`firstfixed.s` is the corrected version:

```
_start:
    mov $60,%rax
    mov $-1,%rdi
    syscall
```

Syscall convention: put the syscall number in `%rax`, put arguments in `%rdi`, `%rsi`, `%rdx`, etc., and then execute `syscall`. For `exit(status)`, the status goes in `%rdi`. So this exits with status -1, and `echo $?` in the shell will confirm it.

add and sub

`firstadd.s` and `firstsub.s` are our first binary operations:

```
_start:
    mov $10,%rbx
    mov $20,%rcx
    add %rbx,%rcx
    mov %rcx,%rdi
    mov $60,%rax
    syscall
```

Syntax note that will bite you all term: in AT&T syntax, `add S, D` means `D = D + S`. The destination is on the *right*. So `add %rbx,%rcx` leaves 30 in `%rcx`. `sub` is analogous: `sub %rbx,%rcx` computes `D - S` and leaves 10 in `%rcx`.

We also did the fun thing of running `gcc -S test.c` on a tiny C file and skimming the output (`test.s`). It's noisy (`.cfi_directives`, frame setup, PLT calls) but the actual work is recognizable: move 10 into a local, load it into `%esi`, load the format string's address via `leaq`, call `printf@PLT`.

4 Lecture 4 — Registers, Calling Convention, Addressing Modes, Control Flow, and “Hello, world”

Today was a dense one: we filled in the register table, introduced the Linux x86-64 calling convention, walked through every addressing mode we'll use for arrays, covered `cmp/jmp`, wrote a for-loop in assembly, and finished by doing `write` and `exit` syscalls by hand.

Logistics

The Linux server admin is in the process of activating accounts for folks who can't log in yet, so hang tight if that's you.

Register reference

The basic x86-64 general-purpose registers and what they were historically used for:

- `%rax` — accumulator (return values, implicit operand for `mul/div`).
- `%rbx` — base register (historically an array base).
- `%rcx` — counter (loops, `rep`).
- `%rdx` — data (I/O, upper half of multiplication/division).
- `%rsi` — source index (string ops).
- `%rdi` — destination index (string ops).
- `%rbp` — base pointer (stack frame base).
- `%rsp` — stack pointer (top of the stack).
- `%r8-%r15` — extra general-purpose registers.
- `%rip` — instruction pointer; we don't write to this directly, but we *do* use it in RIP-relative addressing.

The floating-point/SIMD registers (`%xmm0...`) live in the `assemblyGuide.org` in the repo.

Linux/x86-64 calling convention, first pass

- Return values go in `%rax`.
- The first six integer/pointer arguments are passed in, in order: `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, `%r9`.

Syscalls use a slightly different convention (the fourth argument is `%r10`, not `%rcx`), but for function calls in user code the list above is what you want.

add/sub, again

We re-ran `firstadd.s` from last time as a warm-up; the only new thing is a pair of clarifying comments explaining that `%rdi` is the exit-code argument and 60 in `%rax` selects the `exit` syscall.

Global variables: values, pointers, and `lea`

Three files building on each other:

`addGlobal.s` — define a global with `.quad`, then load, add, store, load-again:

```
.section .data
num: .quad 200

_start:
    mov num,%rbx # load num into %rbx
    add $10,%rbx
    mov %rbx,num # store back
    mov num,%rdi
    mov $60,%rax
    syscall
```


This works, but referencing a symbol like `num` directly makes the assembler emit an absolute address, which is awkward for position-independent code.

`addGlobalBetter.s` — same thing, but using *RIP-relative* addressing:

```
mov num(%rip),%rbx
add $10,%rbx
mov %rbx,num(%rip)
```

Here `num(%rip)` means “the address that is `num`-minus-the-next-instruction bytes away from the current instruction pointer.” The assembler computes that offset for us at build time. This is the form you want in modern code.

`addGlobalLea.s` — now using `lea` to take *the address of* `num`, the way `&` does in C:

```
lea num(%rip),%rbx # rbx = &num
addq $10, (%rbx) # *rbx = *rbx + 10
mov num(%rip),%rdi
mov $60,%rax
syscall
```

The comment block at the bottom of the file spells out the C equivalent:

```
int num = 200;
int* nump = &num;
*nump = *nump + 10;
```

Parentheses around a register mean “dereference.” So `(%rbx)` is “the 8 bytes of memory starting at the address held in `%rbx`.”

From a scalar to an array

Four files, one new addressing idea at a time.

`addArray1.s` — change `num` to an array of four quadwords. The code is unchanged, so it just touches the first element.

`addArray2.s` — to get at element 1, add 8 to the pointer (because one `.quad` is 8 bytes):

```
lea num(%rip),%rbx
addq $8,%rbx
addq $10, (%rbx)
movq (%rbx),%rdi
```

`addArray3.s` — same effect, but using the full “scale-index” addressing mode:

```
lea num(%rip),%rbx
mov $1,%rcx
addq $10, (%rbx,%rcx,8) # *(base + index*8)
movq (%rbx,%rcx,8),%rdi
```

The syntax `(base, index, scale)` computes `base + index*scale`. The scale has to be 1, 2, 4, or 8. This is exactly how C-style array indexing compiles — you’ll see the same form over and over.

`addArray4.s` — same as `addArray2`, but with `.long` (4-byte) elements instead of `.quad`, which forces us to use the `l` suffix (`addl`, `movl`) and the 32-bit register name `%edi`.

Note the subtle point: with `.long`, a step of “one element” is 4 bytes, not 8. In the file as written, we advanced the pointer by 8 bytes, which lands us at element *two* of the `.long` array, not element one. That kind of size/stride mismatch is a classic bug I’ll keep harping on all term.

cmp and jmp

`cmp S, D` does (roughly) $D - S$ and sets flags without saving the result anywhere. The conditional jumps then read those flags:

- `jmp` — unconditional.
- `j1` — $D < S$. (Yes, `cmp S, D` then “if less” compares D to S ; it still surprises me.)
- `jle` — $D \leq S$.
- `je`, `jne` — equal, not equal.
- `jg`, `jge` — greater, greater-or-equal.

The demo, `cmp1.s`:

```
mov $10,%rbx
mov $20,%rcx
mov $-1,%rdi
cmp %rbx,%rcx # flags set as if computing %rcx - %rbx
jge greater
mov $1,%rdi
greater:
mov $60,%rax
syscall
```

Since $20 \geq 10$, we jump to `greater` and exit with status `-1` (or 255 as far as the shell is concerned).

A for-loop in assembly

Two flavors of “sum the numbers from 1 to 10.”

`sumLoop.s` counts up:

```
mov $0,%rbx # accumulator
mov $1,%rcx # counter
loopStart:
add %rcx,%rbx
add $1,%rcx
cmp $10,%rcx
jle loopStart
```

`sumLoopB.s` counts down:

```
mov $0,%rbx
mov $10,%rcx
loopStart:
add %rcx,%rbx
```

```
sub $1,%rcx
cmp $0,%rcx
jg loopStart
```

Both leave 55 in `%rbx`, which we then move into `%rdi` so we can read it out as the exit status.

Buffers, strings, and write

The grand finale, `hello.s`:

```
.section .data
msg: .asciz "Hello world!\n"
hellolen = . - msg

.section .text
.globl _start

_start:
    mov $1,%rax # syscall number for write
    mov $1,%rdi # fd = 1 (stdout)
    lea msg(%rip),%rsi # buffer address
    mov $hellolen,%rdx # number of bytes
    syscall

    xor %rdi,%rdi # exit status 0
    mov $60,%rax # syscall number for exit
    syscall
```

Points of interest:

- `.asciz` lays down the characters followed by a null terminator.
- `hellolen = . - msg` is a classic assembly idiom: the dot is “current address,” so subtracting the start of `msg` gives the length in bytes, including the newline and the terminator. We use that as the third argument to `write`.
- `write` is syscall 1 on Linux x86-64. Arguments are (`fd`, `buf`, `count`) in `%rdi`, `%rsi`, `%rdx`.
- `xor %rdi,%rdi` is the idiomatic way to zero a register — one byte shorter than `mov $0,%rdi`.

Next time we’ll close this out with an `echo`-style program that *reads* from `stdin` and writes back, which means handling buffers properly.

5 Lecture 5 — Hello Again, echo, and the Road from Jumps to Proper Functions

Today was the big one for putting structure around assembly: we re-did “hello, world” carefully, wrote an `echo` that reads and writes a buffer, then walked the whole arc from `jmp`-style subroutines to `call/ret` and finally to real stack frames with local variables. By the end, we had a function in assembly that looks and behaves like a function in C.

Hello, world — this time with feeling

We re-examined `hello.s` slowly, focusing on the pieces we glossed over last time:

```
.section .data
msg: .asciz "Hello world!\n"
hellolen = . - msg

.section .text
.global _start

_start:
    mov $1,%rax # write syscall
    mov $1,%rdi # fd = 1 (stdout)
    lea msg(%rip),%rsi # buffer
    mov $hellolen,%rdx # length
    syscall

    xor %rdi,%rdi
    mov $60,%rax
    syscall
```

A few things worth restating:

- File descriptors in Unix are just integers indexing into a per-process table the kernel maintains. 0 is stdin, 1 is stdout, 2 is stderr, and anything you `open` returns the next free slot. `write` doesn't care what kind of thing the fd refers to — pipe, socket, file, terminal — it all looks the same from up here.
- The `write` syscall *returns* something in `%rax`: the number of bytes it actually wrote (or a negative `errno` on failure). We didn't use it here, but we will.

`helloRet.s` is the tiny variant that does use it:

```
mov $1,%rax
mov $1,%rdi
lea msg(%rip),%rsi
mov $hellolen,%rdx
syscall

mov %rax,%rdi # return value of write becomes our exit status
mov $60,%rax
syscall
```

Run it and then `echo $?` — you'll see the byte count of "Hello world!\n" come out as the exit status. This is the idea I want you to carry forward: `%rax` holds the return value of whatever you just called (syscall or function), and you can feed it into whatever comes next.

Reading input: .bss and the read syscall

To read we need somewhere to read *into*, which is what the `.bss` section is for — uninitialized, zero-filled storage that costs nothing on disk because the loader knows to just hand us zeroed pages:

```
.section .bss
buff: .skip 128
```

`.skip 128` reserves 128 bytes at the label `buff`.

`read` is `syscall 0`, and its arguments mirror `write`:

- `%rdi`: fd (0 = `stdin`)
- `%rsi`: pointer to the buffer
- `%rdx`: maximum number of bytes

And just like `write`, `read` returns the number of bytes actually read in `%rax`.

echo.s: read and then write

Here's `echo.s` end-to-end:

```
.section .bss
buff: .skip 128

.section .text
.global _start

_start:
    mov $0,%rax # read
    mov $0,%rdi # stdin
    lea buff(%rip),%rsi
    mov $128,%rdx
    syscall

    mov %rax,%rdx # bytes-read becomes the count for write
    mov $1,%rax # write
    lea buff(%rip),%rsi
    mov $1,%rdi # stdout
    syscall

    xor %rdi,%rdi
    mov $60,%rax
    syscall
```

The clever step is the first line after `syscall: mov %rax,%rdx`. `read` told us how much it got; we immediately funnel that into the byte-count argument of `write`, so we only echo back what was actually typed. If we hardcoded 128 there instead, we'd dump the uninitialized tail of the buffer too — not a crash, but garbage.

Jump-based “subroutines” — and why they don't scale

Before `call/ret` existed conceptually for us, the hack was to use labels and unconditional jumps. `routine.s` is that hack:

```
rdadder:
    add %rdi,%rdi
    mov %rdi,%rax
    jmp rdret

_start:
```

```

    mov $20,%rdi
    jmp rdadder
rdret:
    mov %rax,%rdi
    mov $60,%rax
    syscall

```

It works: we jump to `rdadder`, it doubles `%rdi`, then `jmp rdret` gets us back. But look at what we had to do — the “return address” is hardcoded into the function body. If we wanted to call `rdadder` from two different places, we’d have to either duplicate it or add some awful switch. It’s not reusable, it’s not encapsulated, it’s not a function.

Real functions: `call` and `ret`

`call` and `ret` are two special instructions that fix the whole thing. Conceptually:

- `call label` pushes the address of the *next* instruction onto the stack and then jumps to `label`.
- `ret` pops that address off the stack and jumps to it.

So `call/ret` are just “save your place, jump, come back” — and because the save uses the stack, you can call as many nested functions as you want.

`fun1.s` is the first real function we wrote:

```

rdadder:
    add %rdi,%rdi
    mov %rdi,%rax
    ret

_start:
    mov $20,%rdi
    call rdadder # %rax = 40
    mov %rax,%rdi
    call rdadder # %rax = 80
    mov %rax,%rdi
    mov $60,%rax
    syscall

```

Now `rdadder` is reusable — we call it twice in a row and compound the doubling. Exit status should be 80.

`fun2.s` is a variant that uses `%r9` as a scratch register instead of clobbering `%rdi`:

```

rdadder:
    mov %rdi,%r9
    add %r9,%r9
    mov %r9,%rax
    ret

_start:
    mov $20,%rdi
    call rdadder # %rdi still 20, %rax = 40

```

```
call rdadder # %rdi still 20, %rax = 40 again
mov %rax,%rdi
mov $60,%rax
syscall
```

Notice what's different from `fun1.s`: the `_start` here calls `rdadder` twice back-to-back without reloading `%rdi` in between. Because `rdadder` no longer modifies `%rdi` (that's the point of using `%r9` as a scratch), both calls see the same input and produce the same output. Exit status is 40, not 80. If you want to compound the effect, you have to feed the previous result back in as the next argument — exactly what `fun1.s` does with its `mov %rax,%rdi` between calls. The lesson: “the function is pure” and “calling it twice doubles the answer” are two very different statements.

Local variables mean we need a stack frame

For anything more interesting than “double a number,” functions need scratch space — local variables. Every process gets a big region of memory called the stack, and every function call gets a contiguous slice of it called a *stack frame*.

Two registers run the show:

- `%rsp` — stack pointer, points at the “top” of the stack.
- `%rbp` — base pointer, points at the “bottom” of the current frame.

Important weirdness: the stack *grows downward*. “Top” is the lowest address. Pushing onto the stack *decreases* `%rsp`; popping *increases* it. A function's local variables live at negative offsets from `%rbp` (e.g. `-8(%rbp)`, `-16(%rbp)`).

`push` and `pop` are the two primitives: `push %rbx` subtracts 8 from `%rsp` and stores `%rbx` there; `pop %rbx` does the inverse.

The canonical function prologue:

1. `push %rbp` — save the caller's base pointer.
2. `mov %rsp,%rbp` — make our new frame start where the stack currently is.
3. `sub $N,%rsp` — carve out N bytes of locals.

And the epilogue undoes it:

1. `mov %rbp,%rsp` — drop all the locals in one move.
2. `pop %rbp` — restore the caller's base pointer.
3. `ret`.

`funStack.s`: a function with locals, by hand

We're computing $f(x) = 20 + 2x$, but deliberately using local variables for both `a = 20` and the intermediate $2x$:

```
fun:
    push %rbp
    mov %rsp,%rbp
    sub $16,%rsp # room for two quadword locals
```

```

    mov %rdi,-8(%rbp) # local1 = x
    add %rdi,-8(%rbp) # local1 += x (so local1 = 2x)
    movq $20,-16(%rbp) # local2 = 20
    movq -16(%rbp),%rax
    add -8(%rbp),%rax # %rax = 20 + 2x

    mov %rbp,%rsp
    pop %rbp
    ret

_start:
    mov $10,%rdi
    call fun
    mov %rax,%rdi
    mov $60,%rax
    syscall

```

With input 10, this exits with status 40. Walk through it once on paper: draw a stack growing downward, watch `push %rbp`, `mov %rsp,%rbp`, and `sub $16,%rsp` carve out the frame. The `-8(%rbp)` and `-16(%rbp)` slots are *your* two local variables; naming them in a comment alongside the code is a lifesaver when this gets bigger.

One subtle detail we’ll formalize later: after `push %rbp` we’ve pushed one 8-byte value, which actually re-aligns `%rsp` to a 16-byte boundary (assuming the caller’s `call` left it 8-aligned, which it does). That matters because the ABI requires `%rsp` to be 16-aligned at the point of a `call`, and some instructions (notably SSE moves) fault if it isn’t.

6 Lecture 6 — Reusing Functions Across Files, Recursion, and Summing an Array

With `call/ret` and stack frames in hand, today was about doing things that actually resemble programming: linking pre-written `readInt/writeInt` helpers in with our code, writing a recursive `factorial`, and walking an array to sum it up.

Logistics

First, grades for the stack assignments are in — nice work overall.

Second, the next coding assignment: write an assembly program that *(i)* declares an array, *(ii)* reads input, *(iii)* does something interesting combining the two, and *(iv)* exits cleanly. Two constraints: you must assemble with `as` and link with `ld` (no `gcc` driver), and you’re encouraged to pull in the `readInt/writeInt` helpers from the assembly guide. Due May 10th.

Third, Discussion 3 is on *posits*, an alternative to IEEE-754. Read `positTutorial.org` in the repo. The short pitch: posits spread their precision more evenly across the number line, so they give up some resolution at very small magnitudes in exchange for much better resolution at moderate ones, and they ditch the menagerie of special cases (no $\pm\infty$, no NaN, just a single NaR). A posit is made up of

- a sign bit,

- *regime* bits (a kind of super-exponent, variable width),
- exponent bits (variable width),
- fraction bits (whatever's left).

The variable widths are what let posits reallocate precision dynamically. We'll leave the details for the tutorial.

Linking in functions from other files

Up to now, everything we've built has been a single `.s` file with a `_start`. Real programs are many files linked together. To use a function defined elsewhere you just need to tell the assembler “this symbol is defined in another file”; the linker will wire it up later.

`readInt.s` defines (and makes global with `.global`) two functions:

```
.section .bss
buffer: .zero 128

.section .text
.global readInt

parseInt:
    movzbl (%rdi,%rcx,1),%r8d # read one byte, zero-extended
    sub $'0',%r8
    imul $10,%rax
    add %r8,%rax
    inc %rcx
    cmp %rsi,%rcx
    jb parseInt
    ret

readInt:
    push %rbx
    mov $0,%rax # sys_read
    mov $0,%rdi # stdin
    lea buffer(%rip),%rsi
    mov $128,%rdx
    syscall

    mov %rax,%rbx # bytes read
    mov $0,%rcx
    mov $0,%rax
    lea buffer(%rip),%rdi
    dec %rbx # trim trailing newline
    mov %rbx,%rsi
    call parseInt
    pop %rbx
    ret
```

Two details worth pointing out. First, `movzbl` (“move byte to long, zero-extended”) loads a single byte of the input buffer into the low 8 bits of `%r8` and clears the upper bits. Without the zero extension, garbage from whatever used `%r8` last would corrupt your arithmetic — classic bug.

Second, `readInt` does `push %rbx / pop %rbx` around its body because `%rbx` is *callee-saved* in the Linux ABI. If we clobbered it, the caller would be entitled to complain.

`writeInt.s` is the mirror image: it takes an integer in `%rdi`, converts it to decimal digits by repeatedly dividing by 10, and prints them with `write`. A few lessons buried in there:

- It allocates a 32-byte buffer *on the stack* and fills it *from the end* — that’s how “print digits in reverse order” naturally handles being emitted left-to-right later.
- It handles negatives by remembering a flag and negating `%rax`, with a special case for `INT_MIN` (-2^{63}), which can’t be negated in two’s complement. The code uses `movabsq` (not plain `movq`) because the constant 2^{63} doesn’t fit in a 32-bit sign-extended immediate.
- It preserves `%rbx` (callee-saved again), and uses `leave` as shorthand for “`mov %rbp,%rsp` then `pop %rbp`.”

To actually call these from another file, we use `.extern` (a declaration, not a definition):

```
.section .text
.global _start
.extern readInt
.extern writeInt

_start:
    call readInt
    mov %rax,%rdi
    call writeInt

    mov $60,%rax
    xor %rdi,%rdi
    syscall
```

That’s `readWriteTest.s`. To build it you assemble each file separately with `as` and then link them all with `ld`:

```
as -o readWriteTest.o readWriteTest.s
as -o readInt.o readInt.s
as -o writeInt.o writeInt.s
ld -o readWriteTest readWriteTest.o readInt.o writeInt.o
```

Type a number, hit enter, and it prints the same number back. Not very exciting on its own, but it’s the scaffolding for everything else today.

Recursion: factorial in assembly

`fact.s` is the first recursive function we’ve written. The idea is familiar — $n! = n \cdot (n - 1)!$, with base case $0! = 1! = 1$ — but implementing it in assembly forces you to be explicit about how the stack carries your intermediate state.

```
fact:
    push %rbp
    mov %rsp,%rbp
    sub $16,%rsp # one local + padding for alignment
```

```

    cmp $1,%rdi
    jle basecase

    mov %rdi,-8(%rbp) # save n before the recursive call
    dec %rdi # %rdi = n - 1
    call fact # recursive call: returns (n-1)! in %rax
    mov -8(%rbp),%rdi # reload our saved n
    imul %rdi,%rax # %rax = n * (n-1)!
    jmp leavefact

basecase:
    mov $1,%rax

leavefact:
    leave
    ret

```

The essential move is `mov %rdi,-8(%rbp)` before the recursive call. `%rdi` is *caller-saved* — any function you call is allowed to clobber it — and recursion means we need `n` again after `fact(n-1)` returns. So we park it in our stack frame. Each recursive activation gets its own `-8(%rbp)` slot, because each activation has its own `%rbp`. That's the whole magic trick.

The `sub $16,%rsp` is worth pausing on. We only need 8 bytes for our one local, but we carve out 16 to keep `%rsp` 16-byte aligned at the next `call`. If this ever drifts you'll see mysterious crashes inside library functions that assume SSE alignment.

The `_start` that drives it:

```

_start:
    mov $1,%rax # print "Enter a number: "
    mov $1,%rdi
    lea message(%rip),%rsi
    mov $16,%rdx
    syscall

    call readInt
    mov %rax,%rdi
    call fact
    mov %rax,%rdi
    call writeInt

    mov $60,%rax
    xor %rdi,%rdi
    syscall

```

Try 5 — you should get 120. Try 20 — 2432902008176640000. Try 21 — you'll overflow a 64-bit signed integer and get nonsense. That silent wraparound is worth remembering.

Summing an array

`sumLoopArray.s` pulls together the array addressing from lecture 4, the counter-loop idiom, and `writeInt`:

```

.section .data
arr: .quad 1,2,3,4

.section .text
.global _start
.extern writeInt

_start:
    mov $0,%rcx # i = 0
    mov $0,%rax # accumulator
    lea arr(%rip),%r15 # base pointer

compare:
    cmp $4,%rcx
    jge done
loop:
    add (%r15,%rcx,8),%rax # %rax += arr[i]
    inc %rcx
    jmp compare
done:
    mov %rax,%rdi
    call writeInt

    mov $60,%rax
    xor %rdi,%rdi
    syscall

```

The key line is `add (%r15,%rcx,8),%rax`: scale-index addressing again, with scale 8 because we’re stepping through `.quad` elements. If you change `arr` to `.long` you must change the scale to 4 *and* switch to the 32-bit forms (`addl, %eax`) — this is the same size/stride trap I warned about in lecture 4, and it will bite people in the assignment.

Output for this program is 10, because $1 + 2 + 3 + 4 = 10$. If you wanted to sum a user-supplied run of numbers instead of a hardcoded array, you already have all the pieces: `readInt` in a loop, accumulate in `%rax`, `writeInt` at the end. That’s a good shape for the upcoming assignment.

7 Lecture 7 — Processes, fork, Threads, and Why Sharing State Hurts

We’re pivoting away from assembly for a bit and back into C, but with a new lens: *concurrency*. Today’s arc was processes vs. threads, `fork` from the bottom up, a quick detour through function pointers because the threading API needs them, and then a first taste of why shared state between threads is a minefield.

Logistics

A few schedule notes worth surfacing: this term, “finals week” lands on what would have been week 10, so I’ve shifted lectures to fit in nine weeks. There’s no in-class final — the last week is buffer time to wrap up the big assignments. Wednesday night of finals week is the hard cutoff for late homework (concretely: before I sit down to grade on Thursday morning).

The next major project — the assembly project from lecture 6 — is due in 13 days, on May 10th. Discussions are still due weekly, quizzes are open and repeatable, and the new course site (still unpublished) is on GitHub — check the pinned announcement.

What is a process?

A *process* is the bundle of context, code, and state that *is* an executing program: its memory, its open file descriptors, its registers, its position in the instruction stream, the works.

In the olden days, a process had *a* thread of execution — exactly one. Multi-user/multi-tasking systems gave each running program a single thread and the OS scheduled between them. The OS *still* schedules: each thread of execution gets a timeslice before being switched out. What’s new (well, new-ish — this has been true for decades now) is that one process can host *multiple* threads of execution, and a real machine has multiple cores that can literally run threads in parallel.

If you want to peek at how Linux represents all this, the relevant header is [include/linux/sched.h](#). The interesting fields on a `task_struct` for our purposes are `parent`, `thread_info`, and `pid`. Linux internally calls both processes and threads “tasks” — the distinction at the kernel level is mostly about what they share with their parent.

fork: the original concurrency primitive

`fork` spawns a new process that is a *clone* of the current one. Same code, same data, same open files, same next instruction. The implications are stranger than they sound:

- The new process picks up at the same line as the parent, because it has the same instruction pointer.
- There’s no shared state between them. Modifying a variable in one does not change it in the other. (Modern kernels save space by having both processes point at the same physical memory pages until somebody writes — “copy-on-write” — but that’s an implementation detail.)
- It’s easy to write `fork`-style code that doesn’t step on itself.
- It’s *harder* to write `fork`-style code that coordinates and shares information, precisely because the state is separate.

`fork1.c` is the bare minimum:

```
printf("We haven't forked yet\n");
fork();
printf("We've now forked\n");
```

Run it and you’ll see “We haven’t forked yet” once and “We’ve now forked” *twice*. After `fork()` returns, there are two processes both sitting at the next line, both about to run that `printf`.

`fork2.c` makes the “no shared state” point concrete:

```
int num = 5;
fork();
num++;
printf("My number is: %d\n", num);
```

Each process has its own copy of `num`, so each one increments its own copy and prints 6. Not 7, not 6 and 7: each process is bookkeeping its own state, full stop. (Aside: a polite program would call `wait()` in the parent so the child can finish before `main` returns. We're skipping that on purpose for now to keep the example short, but it's the kind of thing that bites you in shell pipelines.)

Process IDs and dispatching

If `fork` could only give you two identical clones running in lockstep, it wouldn't be very useful. The trick that makes it useful is that `fork` *returns* different things in the two processes:

- In the parent, `fork()` returns the child's PID (a positive integer).
- In the child, `fork()` returns 0.
- On failure (e.g. out of process slots), it returns -1 in the original process and there is no child.

`fork3.c` just prints the return value:

```
pid_t pid = fork();
printf("The value of pid is: %d\n",pid);
```

You'll see two lines: one with 0 (the child) and one with some big number (the parent's view of the child's PID).

`fork4.c` adds `getpid()`, which always returns "my own PID," so you can see who's who:

```
pid_t pid = fork();
printf("The value of pid is: %d\n",pid);
printf("My name is: %d\n",getpid());
```

The child's two lines will say "pid is 0" and "my name is N," where N is exactly the number the parent printed.

`fork5.c` is the canonical `fork` idiom: *branch on the return value* so parent and child do different things:

```
pid_t pid = fork();
if(pid != 0){
    printf("I'm the parent!\n");
}
else{
    printf("I'm the child!\n");
}
printf("The value of pid is: %d\n",pid);
printf("My name is: %d\n",getpid());
```

This is how every real `fork`-using program is shaped: the parent and the child take different branches and do different work. A shell calls `fork` to make a child, then has the child `exec` the command you typed while the parent waits. A web server forks one child per incoming connection. The pattern is always "one call, two return paths."

For contrast: threads

Threads are the other half of the picture. A thread of execution inside an existing process *shares* state with the other threads in that process: same heap, same globals, same file descriptors. That makes them lighter to spawn (no copy of the whole address space) and makes coordination through shared memory trivial — and it makes interfering with each other equally trivial. We’ll get to that.

The headlines:

- Threads share state, processes don’t.
- Spawning a thread is cheaper than spawning a process.
- Threads need explicit control flow: when you make one, you have to hand it a function to run, and the thread ends when that function returns.

That last point is the wedge for the next topic.

A word on function pointers

To launch a thread you have to give the threading library a *function* to run. C’s way of saying “here, take this function” is a function pointer, which a lot of folks haven’t seen before.

The point of function pointers is that they let you pass a function as a first-class argument. If you’ve used any modern language, you’ve probably used “higher-order functions” that take other functions — `map`, `filter`, `sort-with-a-comparator`, event handlers. C has had this since forever, just with crunchier syntax.

`funpoint1.c` is the warm-up:

```
int doubler(int n){
    return 2*n;
}

void maparr(int arr[], int s,
            int (*func)(int)){
    for(int i=0; i<5; i++){
        arr[i] = func(arr[i]);
    }
}

int main(){
    int arr[5] = {0,1,2,3,4};
    maparr(arr,5,doubler);
    for(int i=0; i<5; i++){
        printf("%d\n",arr[i]);
    }
}
```

The type `int (*func)(int)` reads as “a pointer to a function that takes an `int` and returns an `int`.” The parentheses around `*func` matter — without them you’d be declaring a *function returning a pointer to int*, which is a totally different thing. Inside `maparr`, `func(arr[i])` calls through the pointer; you can also write `(*func)(arr[i])` if you want to be explicit, but the syntax sugar lets you drop the star.

The output is 0 2 4 6 8, because `doubler` got applied to each element in place.

`pthread_create`: actually launching a thread

Now the threading API. The standard on Linux is POSIX threads (`pthread`s). The two calls you'll see today are `pthread_create`, which spawns a thread, and `pthread_join`, which waits for one to finish.

The signature of the thread function is fixed: `void* f(void* p)`. It takes a single `void*` argument and returns a `void*`. Why all the `void*s`? Because this is C's way of doing polymorphism: any pointer can be cast to `void*` and back, so the threading library can hand any function any kind of data without knowing what either of them is.

`thread1.c` is the simplest thing that works:

```
void* threadFun(void* p){
    printf("Hi from a thread!\n");
    return NULL;
}

int main(){
    pthread_t thread;
    pthread_create(&thread, NULL, threadFun, NULL);
    pthread_join(thread, NULL);
    return 0;
}
```

`pthread_create` takes:

1. A `pthread_t*` the library will fill in with the new thread's handle.
2. Optional attributes (`NULL` = defaults).
3. The function to run.
4. The single argument to pass it.

`pthread_join` blocks until the thread exits, and optionally captures its return value (we pass `NULL` because we don't care).

If you forget the join, `main` can return before the thread prints, and depending on timing the program ends without you ever seeing the message. Always join (or detach) your threads.

You'll also notice the `return NULL` at the end of `threadFun`. Returning a pointer when you have nothing to return is C's way of doing an "optional" value: a non-null pointer is your result, `NULL` means "nothing." Modern languages spell this `Maybe` or `Option`; in old-school C it's a pointer convention.

Passing data into a thread

`thread2.c` actually *uses* the argument:

```
void* threadFun(void* p){
    int* arg = (int*)p;
    (*arg)++;
}
```



```

    return NULL;
}

int main(){
    int counter = 0;
    pthread_t thread;
    pthread_create(&thread, NULL, threadFun, &counter);
    pthread_join(thread, NULL);
    printf("Counter is: %d\n", counter);
}

```

We pass `&counter` as the `void*` argument, the thread casts it back to `int*`, increments through it, and the parent sees `Counter is: 1` after the join. Crucially, this works because threads share memory: the thread's `*arg` and `main`'s `counter` are literally the same bytes. Try this with `fork` and you'd see no change in the parent — that's the whole difference.

The danger of shared state

Now the punchline of the lecture. `thread3.c` launches 100 threads, each of which increments `counter` 10,000 times:

```

void* threadFun(void* p){
    int* arg = (int*)p;
    int temp = *arg;
    for(int i=0; i<10000; i++){
        temp++;
    }
    *arg = temp;
    return NULL;
}

int main(){
    int counter = 0;
    pthread_t thread[100];
    for(int i=0; i<100; i++){
        pthread_create(&thread[i], NULL, threadFun, &counter);
    }
    for(int i=0; i<100; i++){
        pthread_join(thread[i], NULL);
    }
    printf("Counter is: %d\n", counter);
}

```

You'd hope, naively, for $100 \cdot 10,000 = 1,000,000$. You will not get 1,000,000. Run it a few times and you'll get different numbers, all far short.

The bug is the read-modify-write pattern in the body. Each thread:

1. reads `*arg` into a local `temp`,
2. does its 10,000 increments on `temp`,
3. writes `temp` back to `*arg`.

While thread A is doing its 10,000 increments, threads B, C, D, ... are also reading `*arg`, doing their increments, and writing back. Whoever writes last wins; everyone else's work is overwritten. This is a *race condition*, and it is the canonical failure mode of shared-state concurrency.

`thread4.c` is the same code with one tiny addition — a `usleep(100)` inside the inner loop, and only 100 iterations per thread instead of 10,000:

```
for(int i=0; i<100; i++){
    temp++;
    usleep(100);
}
```

The `usleep` forces the OS to interleave the threads more visibly — it's much harder for one thread to “win” the race by finishing its loop in a single timeslice. The result: `counter` ends up close to 100, not 10,000. Each of the 100 threads is essentially clobbering everyone else's writes.

The moral isn't “don't use threads.” The moral is that *any time multiple threads touch the same memory, you need some form of synchronization* — mutexes, atomics, message passing, something. We covered the symptom today. The next chunk of the term is going to be about the cures.

8 Lecture 8 — Returning From Threads, Mutexes, and Critical Sections

Picking up where last lecture left us: we know how to spawn threads, we know how to pass them an argument, and we know that naïvely sharing state between them is a recipe for lost updates. Today's arc is the natural follow-up: how do you get a value *out* of a thread, what's the gotcha when you want to give a *group* of threads each their own argument, and what is this “mutex” thing that's supposed to fix the race condition? And then, just as importantly: when does fixing the race condition with a mutex *undo* the whole point of going concurrent in the first place?

Returning a value from a thread

We've passed data *into* threads via the `void*` argument. The flip side is the return value. Recall that the thread function's signature is `void* f(void* p)` — it returns a `void*`, and `pthread_join` can capture that pointer for you if you hand it a `void**`.

`threadRet.c` is the right way to do it:

```
void* worstIncrement(void* i){
    // 1. make space for the return value (malloc)
    // 2. put the value to be returned in it
    // 3. return the pointer
    int* p = malloc(sizeof(int));
    *p = *((int*)i)+1;
    return p;
}

int main(){
    pthread_t thread;
    int argVal = 10;
    void* retVal;
    pthread_create(&thread, NULL, worstIncrement, &argVal);
}
```

```
pthread_join(thread,&retVal);
printf("The returned value was: %d\n",*((int*)retVal));
return 0;
}
```

The pattern is: `malloc` a fresh chunk of memory inside the thread, store the value there, and return the pointer. The caller joins, gets the pointer back via `retVal`, and can dereference it because that memory lives on the heap and outlives the thread that allocated it. (Yes, in real code you’d `free` it afterwards. We’re keeping the example tight.)

Now the “why `malloc`” part. `threadRetBad.c` is the exact same idea but with the value living on the stack instead:

```
void* worstIncrement(void* i){
    int p; // lives in this stack frame, e.g. -4(%rbp)
    p = *((int*)i)+1;
    return &p; // returns a pointer into the about-to-die frame
}
```

The function returns a pointer to a local variable. The instant the function returns, that stack frame is conceptually gone — whatever the next function call does will reuse those bytes. By the time `main` dereferences `retVal` inside its `printf`, that address is pointing at junk. (Or, on a bad day, at memory the process isn’t allowed to read, and you segfault.)

The whole reason we worked with the assembly view of stack frames earlier in the term is so this kind of bug stops being mysterious. “Returning the address of a stack local” isn’t a vague hazard — it is concretely returning something like `-4(%rbp)` after `%rbp` has been restored to somebody else’s frame.

A caveat on passing arguments to a group of threads

Last time we saw a single thread happily share an integer with `main`. Things get subtler when you spin up *many* threads and want each one to get its own distinct argument.

`threadArgs.c` is the right way:

```
void* sayHi(void* p){
    int id = *((int*)p);
    printf("Hi my name is: %d\n",id);
    return NULL;
}

int main(){
    pthread_t threads[20];
    int threadArgs[20];
    for(int i=0; i<20; i++){
        threadArgs[i] = i;
    }
    for(int i=0; i<20; i++){
        pthread_create(&threads[i],NULL,sayHi,threadArgs+i);
    }
    for(int i=0; i<20; i++){
        pthread_join(threads[i],NULL);
    }
}
```

```

}
return 0;
}

```

Each thread gets the address of a *different* slot in the `threadArgs` array, and those slots have already been initialized to 0, 1, 2, ..., 19. Each thread reads its own slot. Easy.

`threadArgsBad.c` is what you'd write if you weren't paying attention — pass each thread the address of the loop counter itself:

```

for(int i=0; i<20; i++){
    pthread_create(&threads[i], NULL, sayHi, &i);
}

```

Now every single thread is given the same pointer: `&i`. Because the loop is racing ahead while the threads are spinning up, each thread reads `*p` *whenever it happens to get scheduled*, which means by the time it actually runs, `i` may have already advanced. You'll see duplicate IDs, missing IDs, IDs of 20 (the post-loop value), all sorts of garbage. Worse, after `main` leaves the loop, `i` goes out of scope at some point, and now your threads are reading freed stack memory — the same hazard as the previous section, dressed up differently.

The fix and the bug both come from the same observation: in C, `&i` is the address of *the variable*, not a snapshot of its current value. If you want each thread to see a distinct value, you need a distinct *location* per thread. That's what the `threadArgs` array buys you.

Mutexes: the simplest fix for the race

Recall `thread3.c/thread4.c` from last lecture: 100 threads each doing a read-modify-write against a shared `counter`, and the final answer being absolutely nowhere near 100·100. `threadCounterBad.c` in this lecture's directory is the same idea, kept around for comparison:

```

void* threadFun(void* p){
    int* arg = (int*)p;
    int temp = *arg;
    for(int i=0; i<100; i++){
        temp++;
        usleep(100);
    }
    *arg = temp;
    return NULL;
}

```

Each thread reads `*arg`, increments a local copy 100 times (with `usleep` in there to make sure the OS interleaves aggressively), and writes the local copy back. They clobber each other.

The mechanical fix is a *mutex* — short for “mutual exclusion.” A mutex is a little flag managed by the OS that exactly one thread at a time can “hold.” If you try to lock a mutex that somebody else is already holding, your thread sleeps until the holder unlocks it.

`threadCounterMutex.c` bolts a mutex onto the bad version:

```

pthread_mutex_t m;

```

```

void* threadFun(void* p){
    int* arg = (int*)p;
    pthread_mutex_lock(&m);
    // first thread here runs through; everyone else waits
    // until somebody calls pthread_mutex_unlock
    int temp = *arg;
    for(int i=0; i<100; i++){
        temp++;
        //usleep(100);
    }
    *arg = temp;
    pthread_mutex_unlock(&m);
    return NULL;
}

int main(){
    pthread_mutex_init(&m, NULL);
    int counter = 0;
    pthread_t thread[100];
    for(int i=0; i<100; i++){
        pthread_create(&thread[i], NULL, threadFun, &counter);
    }
    for(int i=0; i<100; i++){
        pthread_join(thread[i], NULL);
    }
    printf("Counter is: %d\n", counter);
    pthread_mutex_destroy(&m);
    return 0;
}

```

The pattern to internalize:

- `pthread_mutex_init` once, before any threads see it.
- `pthread_mutex_lock` at the top of the protected region. This blocks if somebody else holds it.
- `pthread_mutex_unlock` at the bottom. *Every* path out of the protected region needs to unlock, or the next thread waits forever.
- `pthread_mutex_destroy` once, after all the threads have joined.

With the lock in place, the read-increment-write block becomes indivisible from any other thread's point of view. The final `counter` is now exactly $100 \cdot 100 = 10,000$, every single run.

`threadCounterMutex2.c` is the same code but with the `usleep(100)` re-enabled inside the loop. The point is to show that the answer is *still correct* — you just have to wait a lot longer for it. We'll come back to that in a minute.

Bundling the lock with the data

Putting the mutex in a global variable works, but in any nontrivial program you'd rather keep the lock and the data it protects together. `threadCounterStruct.c` does exactly that:

```

struct ThreadState {

```

```

pthread_mutex_t mut;
int counter;
};

void* threadFun(void* p){
    struct ThreadState* arg = (struct ThreadState*)p;
    pthread_mutex_lock(&arg->mut);
    int temp = arg->counter;
    for(int i=0; i<100; i++){
        temp++;
        usleep(100);
    }
    arg->counter = temp;
    pthread_mutex_unlock(&arg->mut);
    return NULL;
}

int main(){
    struct ThreadState s;
    s.counter = 0;
    pthread_mutex_init(&s.mut, NULL);

    pthread_t thread[100];
    for(int i=0; i<100; i++){
        pthread_create(&thread[i], NULL, threadFun, &s);
    }
    for(int i=0; i<100; i++){
        pthread_join(thread[i], NULL);
    }

    printf("Counter is: %d\n", s.counter);
    pthread_mutex_destroy(&s.mut);
    return 0;
}

```

Same answer, same correctness. The win is that the mutex and the counter it protects are now *one object*: if you pass the struct to another function, the lock goes with it. There’s no globally visible `m` for somebody else’s code to accidentally lock “the same” mutex against unrelated data. This is the shape you’ll want for assignment 3.

The problem with overusing mutexes

Here’s the trap. Wrapping more code in a mutex makes more bugs go away — so why not just wrap *everything*?

Look at `threadCounterMutex2.c` again. The lock is held across the entire 100-iteration loop, including the `usleep(100)` inside it. That means while thread A is sitting inside its loop (sleeping!), threads B, C, D, . . . , all 99 others, are blocked on `pthread_mutex_lock`, doing nothing. Each thread runs in strict sequence after the last one finishes.

If using mutexes too aggressively means each thread runs to completion before the next one starts, you’ve ended up with synchronous code with extra steps. The threads aren’t buying you anything

in terms of efficiency — if anything, things are *worse*, because you’ve got the cost of all those context switches and thread creations on top of a workload that is effectively a for loop.

That’s pretty bad! And it leads directly to the next idea.

Critical sections

The *critical section* is the part of the code that runs between a lock and its corresponding unlock — the region only one thread can be inside at a time.

A big part of designing concurrent code well is figuring out what should and what *should not* be in the critical section. Things that have to be in there: any read-modify-write of shared state, any multi-step update where another thread reading the half-updated state would be wrong. Things that should *not* be in there: actual work. Number-crunching on local variables. I/O that doesn’t touch shared state. Anything that can be done without seeing the shared data.

The mental rule of thumb: a critical section should be as small as possible while still being correct. You hold the lock just long enough to safely grab the next chunk of work or to commit the result of your chunk, and you do the chunk itself with the lock *not* held.

How this connects to assignment 3

Assignment 3 is going to ask you to process a very, very large array — think millions to hundreds of millions of elements — in parallel with a fixed pool of thread workers. If each thread does *all* of its work inside a critical section, you’ve just rebuilt `threadCounterMutex2.c` at scale, and it’ll be ridiculously slow.

The trick is to figure out what actually needs to be guarded. The big hint: if you make sure each thread gets its *next chunk of work* assigned under a lock — so two threads never claim the same chunk — then the actual processing of that chunk can happen with no lock held at all. Locking the bookkeeping is fast; locking the work is not.